

# Assignment I

Git repository link:

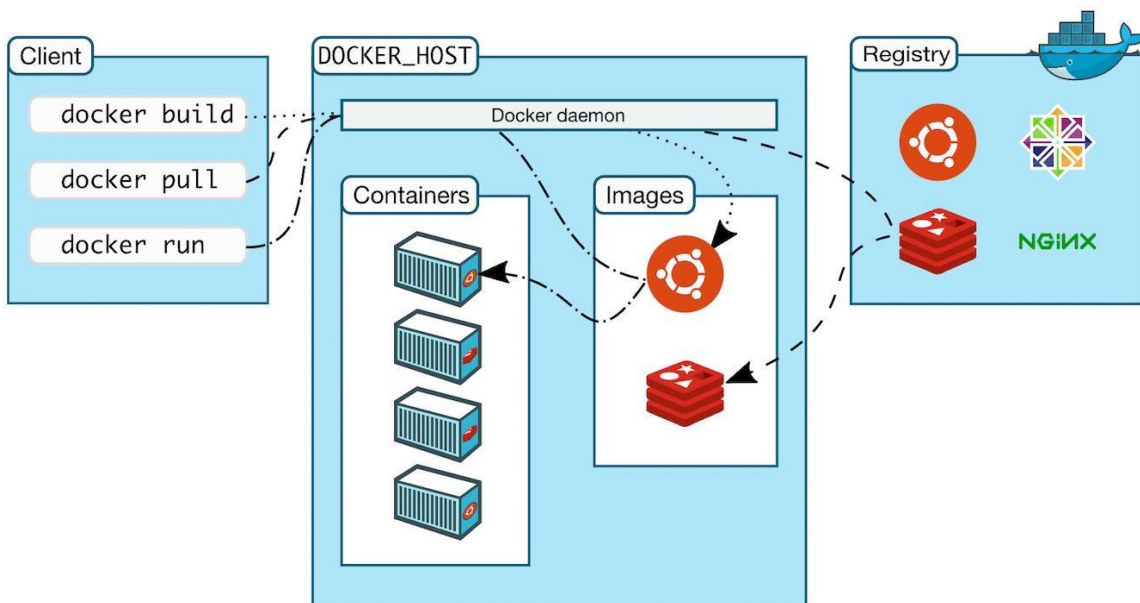
[https://github.com/jucollas/operating\\_systems.git](https://github.com/jucollas/operating_systems.git)

## Activity No. 5: Docker practical session [10%]

Docker is a command line-based software allowing users to manipulate images and create application containers.

As presented in the course, Docker consists of two elements:

- a client, to receive commands from the user
- a server, to execute commands and manage images and containers



### Docker commands architecture

Typing this command will give the Client and Server versions available on your computer. **Paste a screenshot of result**

`docker version`

```
> docker version
Client:
Version:      27.0.3
API version:  1.46
Go version:   go1.21.11
Git commit:   7d4bcd8
Built:        Sat Jun 29 00:03:32 2024
OS/Arch:      windows/amd64
Context:      desktop-linux

Server: Docker Desktop 4.32.0 (157355)
Engine:
Version:      27.0.3
API version:  1.46 (minimum version 1.24)
Go version:   go1.21.11
Git commit:   662f78c
Built:        Sat Jun 29 00:02:50 2024
OS/Arch:      linux/amd64
Experimental: false
containerd:
Version:      1.7.18
GitCommit:    ae71819c4f5e67bb4d5ae76a6b735f29cc25774e
runc:
Version:      1.7.18
GitCommit:    v1.1.13-0-g58aa920
docker-init:
Version:      0.19.0
GitCommit:    de40ad0
```

Usage, options and a full list of available commands can be accessed through the command line in a terminal. Type the following command

```
docker --help
```

The general usage of a Docker command line is as follows:

```
docker [OPTIONS] COMMAND [arg...]
```

## Questions

1. How many arguments are absolutely required by the command 'docker pull' ?  
R// El comando docker pull requiere **un argumento** absolutamente necesario, que es el nombre de la **imagen** que deseas descargar, a esto se le puede adicionar la etiqueta (tag) y el identificador (digest).

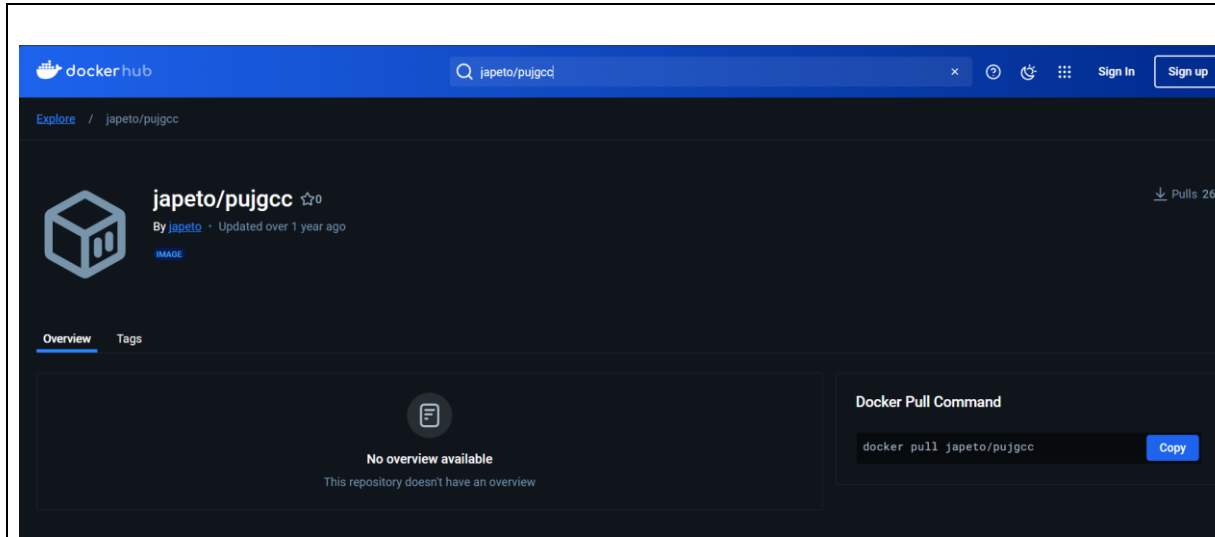
```
docker pull [OPTIONS] NAME[:TAG|@DIGEST]
```

2. Do you remember what a registry is?

R// Un registro (**registry**) en Docker es un servicio donde se almacenan y distribuyen las imágenes de Docker.

### Download a predefined image available on the DockerHub

In a web browser, navigate to the DockerHub : <https://hub.docker.com/> In the top search bar, type : japeto/pujgcc and paste a screenshot



Our course has images available for the development of practical sessions.

### Questions

1. How many times was the *japeto* image downloaded ?

R// 262

Execute the command inside a terminal.

**Paste a screenshot of result**

```
docker pull japeto/pujgcc:v0.12
```

```
> docker pull japeto/pujgcc:v0.12
```

```
What's next:
```

```
View a summary of image vulnerabilities and recommendations → docker scout quickview ja  
eto/pujgcc:v0.12
```

```
Error response from daemon: manifest for japeto/pujgcc:v0.12 not found: manifest unknown: m  
anifest unknown
```

You will get an error as this image has no default tag (“latest”). So we need to specify one in the command line.

Go to the “Tags” tab and copy the pull command of version latest

### **Paste a screenshot of result**

```
docker pull japeto/pujgcc:latest
```

```
> docker pull japeto/pujgcc:latest
latest: Pulling from japeto/pujgcc
Digest: sha256:9b3d7bbf410396d0a26e6bcffbb54e4239736d5a23ad694b03d93c26f9fc6868
Status: Image is up to date for japeto/pujgcc:latest
docker.io/japeto/pujgcc:latest

What's next:
  View a summary of image vulnerabilities and recommendations → docker scout quickview japeto/pujgcc:latest
```

#### Question:

1. How many times do you see 'Pull complete' displayed ? Why ?  
R// Cuando Docker descarga una imagen, a menudo está compuesta por varias capas. Cada capa representa una parte del sistema de archivos que ha sido creada o modificada. El proceso de descarga de una imagen implica obtener cada una de estas capas desde el registro. No pudimos determinar cuántas veces aparece "pull complete" ya que se había descargado con anterioridad la imagen.

Now, to be sure that the image was correctly pulled, let's see the list of all available downloaded images inside our workspace. **Paste a screenshot of result**

```
docker image ls
```

```
> docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
docker/welcome-to-docker	latest	c1f619b6477e	10 months ago	18.6MiB
japeto/pujgcc	latest	05141310a94c	19 months ago	1.39GiB

#### Question:

1. What is the size of the japeto/pujgcc image ?  
R// el tamaño de imagen japeto/pujgcc es de 1.39GB

### **Perform a task using a pulled image**

Among the Docker commands, we will now use the 'run' command.

#### Question:

1. What are the options and parameters of the 'run' command ?  
R// Los parámetros son:  
"docker run [OPTIONS] IMAGE [COMMAND] [ARG...]"

**docker run:** Comando principal para crear y ejecutar un contenedor.

**[OPTIONS]:** Opciones adicionales que puedes proporcionar para configurar el contenedor. Estas opciones pueden incluir configuraciones como puertos expuestos, variables de entorno, volúmenes montados, entre otros. Estas son todas las opciones:

```
Options:
  --add-host list          Add a custom host-to-IP mapping
                           (host:ip)
  --annotation map        Add an annotation to the
                           container (passed through to the
                           OCI runtime) (default map[])
-a, --attach list         Attach to STDIN, STDOUT or STDERR
--blkio-weight uint16     Block IO (relative weight),
                           between 10 and 1000, or 0 to
                           disable (default 0)
--blkio-weight-device list Block IO weight (relative device
                           weight) (default [])
--cap-add list            Add Linux capabilities
--cap-drop list          Drop Linux capabilities
--cgroup-parent string    Optional parent cgroup for the
                           container
--cgroupns string         Cgroup namespace to use
                           (host|private)
                           'host': Run the container in
                           the Docker host's cgroup
                           namespace
                           'private': Run the container in
                           its own private cgroup namespace
                           '': Use the cgroup
                           namespace as configured by the
                           default-cgroupns-mode
                           option on the daemon (default)
--cidfile string          Write the container ID to the file
--cpu-period int          Limit CPU CFS (Completely Fair
                           Scheduler) period
--cpu-quota int           Limit CPU CFS (Completely Fair
                           Scheduler) quota
--cpu-rt-period int       Limit CPU real-time period in
                           microseconds
--cpu-rt-runtime int      Limit CPU real-time runtime in
                           microseconds
-c, --cpu-shares int       CPU shares (relative weight)
--cpus decimal            Number of CPUs
--cpuset-cpus string       CPUs in which to allow execution
                           (0-3, 0,1)
--cpuset-mems string       MEMs in which to allow execution
                           (0-3, 0,1)
-d, --detach              Run container in background and
                           print container ID
--detach-keys string       Override the key sequence for
                           detaching a container
--device list             Add a host device to the container
--device-cgroup-rule list Add a rule to the cgroup allowed
                           devices list
--device-read-bps list    Limit read rate (bytes per
                           second) from a device (default [])
--device-read-iops list   Limit read rate (IO per second)
                           from a device (default [])
--device-write-bps list   Limit write rate (bytes per
                           second) to a device (default [])
--device-write-iops list  Limit write rate (IO per second)
                           to a device (default [])
--disable-content-trust   Skip image verification (default
                           true)
--dns list                Set custom DNS servers
--dns-option list         Set DNS options
--dns-search list         Set custom DNS search domains
--domainname string       Container NIS domain name
--entrypoint string       Overwrite the default ENTRYPOINT
                           of the image
-e, --env list            Set environment variables
```

<code>--env-file list</code>	Read in a file of environment variables
<code>--expose list</code>	Expose a port or a range of ports
<code>--gpu gpu-request</code>	GPU devices to add to the container ('all' to pass all GPUs)
<code>--group-add list</code>	Add additional groups to join
<code>--health-cmd string</code>	Command to run to check health
<code>--health-interval duration</code>	Time between running the check (ms s m h) (default 0s)
<code>--health-retries int</code>	Consecutive failures needed to report unhealthy
<code>--health-start-interval duration</code>	Time between running the check during the start period (ms s m h) (default 0s)
<code>--health-start-period duration</code>	Start period for the container to initialize before starting health-retries countdown (ms s m h) (default 0s)
<code>--health-timeout duration</code>	Maximum time to allow one check to run (ms s m h) (default 0s)
<code>--help</code>	Print usage
<code>-h, --hostname string</code>	Container host name
<code>--init</code>	Run an init inside the container that forwards signals and reaps processes
<code>-i, --interactive</code>	Keep STDIN open even if not attached
<code>--ip string</code>	IPv4 address (e.g., 172.30.100.104)
<code>--ip6 string</code>	IPv6 address (e.g., 2001:db8::33)
<code>--ipc string</code>	IPC mode to use
<code>--isolation string</code>	Container isolation technology
<code>--kernel-memory bytes</code>	Kernel memory limit
<code>-l, --label list</code>	Set meta data on a container
<code>--label-file list</code>	Read in a line delimited file of labels
<code>--link list</code>	Add link to another container
<code>--link-local-ip list</code>	Container IPv4/IPv6 link-local addresses
<code>--log-driver string</code>	Logging driver for the container
<code>--log-opt list</code>	Log driver options
<code>--mac-address string</code>	Container MAC address (e.g., 92:d0:c6:0a:29:33)
<code>-m, --memory bytes</code>	Memory limit
<code>--memory-reservation bytes</code>	Memory soft limit
<code>--memory-swap bytes</code>	Swap limit equal to memory plus swap: '-1' to enable unlimited swap
<code>--memory-swappiness int</code>	Tune container memory swappiness (0 to 100) (default -1)
<code>--mount mount</code>	Attach a filesystem mount to the container
<code>--name string</code>	Assign a name to the container
<code>--network network</code>	Connect a container to a network
<code>--network-alias list</code>	Add network-scoped alias for the container
<code>--no-healthcheck</code>	Disable any container-specified HEALTHCHECK
<code>--oom-kill-disable</code>	Disable OOM Killer
<code>--oom-score-adj int</code>	Tune host's OOM preferences (-1000 to 1000)
<code>--pid string</code>	PID namespace to use
<code>--pids-limit int</code>	Tune container pids limit (set -1 for unlimited)
<code>--platform string</code>	Set platform if server is multi-platform capable
<code>--privileged</code>	Give extended privileges to this container
<code>-p, --publish list</code>	Publish a container's port(s) to the host
<code>-P, --publish-all</code>	Publish all exposed ports to random ports
<code>--pull string</code>	Pull image before running ("always", "missing", "never") (default "missing")

```

("always", "missing", "never")
(default "missing")
-q, --quiet          Suppress the pull output
--read-only          Mount the container's root
                    filesystem as read only
--restart string      Restart policy to apply when a
                    container exits (default "no")
--rm                Automatically remove the
                    container and its associated
                    anonymous volumes when it exits
--runtime string      Runtime to use for this container
--security-opt list   Security Options
--shm-size bytes      Size of /dev/shm
--sig-proxy           Proxy received signals to the
                    process (default true)
--stop-signal string Signal to stop the container
--stop-timeout int    Timeout (in seconds) to stop a
                    container
--storage-opt list    Storage driver options for the
                    container
--sysctl map          Sysctl options (default map[])
--tmpfs list          Mount a tmpfs directory
-t, --tty            Allocate a pseudo-TTY
--ulimit ulimit       Ulimit options (default [])
-u, --user string     Username or UID (format:
                    <name|uid>[:<group|gid>])
--usersns string      User namespace to use
--uts string          UTS namespace to use
-v, --volume list     Bind mount a volume
--volume-driver string Optional volume driver for the
                    container
--volumes-from list   Mount volumes from the specified
                    container(s)
-w, --workdir string  Working directory inside the
                    container

```

**IMAGE:** El nombre de la imagen de Docker que se utilizará para crear el contenedor. Puedes especificar una etiqueta opcional para usar una versión específica de la imagen.

**[COMMAND]:** (Opcional) El comando que deseas ejecutar en el contenedor. Si no se especifica, Docker ejecutará el comando predeterminado definido en la imagen.

**[ARG...]:** (Opcional) Argumentos que se pasan al comando especificado.

```
docker run --help
```

As displayed in the terminal, the description of the command is 'Run a command in a new container'.

Question :

1. What is the difference between an image and a container ?

R// En Docker, una imagen es una plantilla estática e inmutable que contiene el código, las dependencias y el entorno necesario para ejecutar una aplicación. Se utiliza como un plano para crear contenedores. Un contenedor, por otro lado, es una instancia en ejecución de una imagen que es ligera y mutable. Los contenedores usan la imagen como base para ejecutar la aplicación y pueden ser iniciados, detenidos y modificados durante su ciclo de

vida. Las imágenes definen el entorno y las configuraciones, mientras que los contenedores son las instancias activas que ejecutan el software basado en esas imágenes.

Now, to run the application, execute the following command:

**Paste a screenshot of result**

```
docker run japeto/pujgcc bash --help
```

```
> docker run japeto/pujgcc bash --help
GNU bash, version 4.3.30(1)-release-(x86_64-pc-linux-gnu)
Usage: bash [GNU long option] [option] ...
        bash [GNU long option] [option] script-file ...
GNU long options:
    --debug
    --debugger
    --dump-po-strings
    --dump-strings
    --help
    --init-file
    --login
    --noediting
    --noprofile
    --norc
    --posix
    --rcfile
    --restricted
    --verbose
    --version
Shell options:
    -ilrsD or -c command or -O shopt_option          (invocation only)
    -abefhkmnptuvxBCHP or -o option
Type `bash -c "help set"' for more information about shell options.
Type `bash -c help' for more information about shell builtin commands.
Use the `bashbug' command to report bugs.
```

**Congratulations!**

**You just successfully downloaded and used your first Docker image!**

Running *PUJGCC* without parameters was interesting as a demonstration of Docker's features. But if we want to really run *PUJGCC*, we also need to provide parameters and, most importantly, input files.

### Find the paths to bind

To bind our current folder to the `/data/` folder located inside a container, we first need the absolute path of the current folder, obtained through the unix `pwd` command. **Paste a screenshot of result**

```
pwd
```



```
> pwd  
Path  
C:\Users\PC
```

This path will be used in further commands through `${PWD}`.

Instead of running `ls` command to `/home/` files, we will now just list the content of the `/data/` folder inside the container but bind with the host.

**Paste a screenshot of result**

```
docker run japeto/pujgcc:latest ls /data
```

```
> docker run japeto/pujgcc:latest ls /data  
ls: cannot access /data: No such file or directory
```

If nothing appears, it is normal: the folder is empty and only serves as a “*branching point*”.

We now have the paths of the two folders we want to bind together.

**Bind a local folder into a container**

To perform the folder mapping between the current folder and `/data` inside the image, the syntax is simple. **Paste a screenshot of result**

```
docker run -v ${PWD}:/data/ japeto/pujgcc:latest ls /data/
```

```
> docker run -v ${PWD}:/data/ japeto/pujgcc:latest ls /data/  
helloworld.c
```

Question:

1. Is the displayed list the same as what is in your current folder?  
R// Si, las carpetas tiene los mismos archivos.

Finally, we can run C on a C or C++ file located in the Data folder. Change the name of the file to any of the provided files.

**Paste a screenshot of result**

```
docker run -v ${PWD}:/data/ japeto/pujgcc:latest gcc /data/helloworld.c
```

```
> docker run -v ${PWD}:/data/ japeto/pujgcc:latest gcc /data/helloworld.c  
Collazos ~\data 1.048s
```

## Restart and detach a container

Learn how to re-use a container where you installed something

Use the start command to restart the container created in the last exercise **Paste a screenshot of result**

```
docker start mycontainer
```

```
> docker start mycontainer  
mycontainer
```

Go back to the container using the exec command instead of the run command. **Paste a screenshot of result**

```
docker exec -ti mycontainer /bin/bash
```

```
> docker exec -ti mycontainer /bin/bash  
root@772be3f098fc:/# ls  
bin    dev    home  lib64  mnt    proc   run    srv    tmp    var  
boot   etc    lib   media  opt    root   sbin   sys    usr  
root@772be3f098fc:/#
```

Question :

1. What happens now when you exit the container? Is it stopped?  
R// No se detiene, sigue ejecutándose el contenedor.

Check with and **Paste a screenshot of result**

```
docker ps -l
```

```
> docker ps -l  
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS        PORTS   NAMES  
772be3f098fc   japeto/pujgcc  "/bin/bash"             14 minutes ago Up 2 minutes  mycontainer
```

In fact, the container keeps on running. This is because re-starting a container turns it into a “detached process” running in the background. Alternatively, we could have added the -d option to the first docker run command, creating directly a detached container.

Finally, you can stop the container.

```
docker stop mycontainer
```

```
> docker stop mycontainer  
mycontainer
```